

# RL- models evaluation on Snake

Assesment of DQN, DDQN algorithms


# Problem statement

- Environment:
  - 2 dimensional space, bounded. A snake moves around this space looking for food to eat.  
When the snake eats it grows in length. Do not hit the barrier or the snake itself!
- Task:
  - Train an agent with DQ Learning to achieve the best possible score
- Assess improvements:
  - Use DQN extension
  - Inject some experiences before the training process

# Tools




# Agent interactions



# Environment


- Game state vector:
  - Danger
    - Straight, left, right
  - Current moving directions
    - up, down, left, right
  - Current food direction
    - up, down, left, right

example : [0 1 0 1 0 0 0 1 0 0 1]



# Percepts


- Reward
  - If collision -> -10
  - If chop food -> +10
  - else 0
- Done:
  - Whether the game is finished or not



# Actions

- 3 possible actions
  - Forward, turn-left, turn-right

example : [0 0 1]



# Training process - Optimal Bellman equation

$$Q(s, a) = \mathbb{E} \left[ \sum_{t=0}^{\infty} \gamma^t R_t \mid s, a \right]$$


$$Q(s, a) = \sum_{s', r} p(s', r | s, a) \left[ r + \gamma \max_{a'} Q(s', a') \right]$$

# Q-learning - Base approach

$$\text{New } Q(s, a) = Q(s, a) + \alpha [R(s, a) + \gamma \max Q'(s', a') - Q(s, a)]$$


TD-error

- New Q Value for that state and the action
- Learning Rate
- Reward for taking that action at that state
- Current Q Values
- Maximum expected future reward given the new state ( $s'$ ) and all possible actions at that new state.
- Discount Rate



# Deep learning approach- Deep Q-Learning

- Deep Q-Network
  - In DQN, the Q-function is represented implicitly by a neural network



# Model - DQN

---

**Algorithm 1** Deep Q-Network (DQN)

Initialize replay memory  $D$  with capacity  $N$

Initialize Q-network with random weights

**for** episode in range(total\_episodes) **do**

    Initialize state  $s$

**for** step in range(max\_steps\_per\_episode) **do**

        Choose action  $a$  using epsilon-greedy policy based on Q-network

        Take action  $a$ , observe reward  $r$  and next state  $s'$

        Store  $(s, a, r, s')$  in replay memory  $D$

        Sample a random minibatch of experiences  $(s, a, r, s')$  from  $D$

        Calculate target Q-values:

**if**  $s'$  is terminal state **then**

$target = r$

**else**

$target = r + \gamma \max_{a'}(Q(s', a'))$

**end if**

        Update the Q-network weights using gradient descent:

$L(\theta) = \text{MSE}(Q(s, a), target)$

$\theta \leftarrow \theta - \alpha \nabla_{\theta} L(\theta)$

        Backpropagate the loss and update the Q-network weights


        Update state  $s = s'$

**end for**

    Reduce epsilon (exploration rate) over time (e.g.,  $\epsilon^* = \epsilon_{\text{decay}}$ )


**end for**

# Model - DQN



# Models - Policy and action selection strategies


- Like QL, this is an off-policy algorithm
  - The policy we try to learn is not the same as the policy used for action selection during training.  
Agent learns from its past experiences, regardless of the policy used to generate those experiences.
- Epsilon-greedy
  - Trade-off between exploration and exploitation
  - epsilon is reduced while the training goes on



# Models

- DQN extension: Double Deep Q-Network
  - This method handles the problem of the overestimation of Q-values.  
During the training, especially at the beginning, the Q network estimates which is the best action to take, but possibly the actual Q value is overestimated
  - When we compute the Q target, we use two networks to decouple the action selection (online network) from the target Q value generation (target network).

# Models – DDQN



# Models – DDQN

---

**Algorithm 2** Double Deep Q-Network (Double DQN)


---

Initialize replay memory  $D$  with capacity  $N$   
Initialize Q-network  $Q$  and target Q-network  $Q'$  with random weights

**for** episode in range(total\_episodes) **do**  
    Initialize state  $s$   
    **for** step in range(max\_steps\_per\_episode) **do**  
        Choose action  $a$  using epsilon-greedy policy based on Q-network  $Q$   
        Take action  $a$ , observe reward  $r$  and next state  $s'$   
        Store  $(s, a, r, s')$  in replay memory  $D$   
        Sample a random minibatch of experiences  $(s, a, r, s')$  from  $D$   
  
        Calculate target Q-values using target network  $Q'$ :  
        **if**  $s'$  is terminal state **then**  
             $target = r$   
        **else**  
             $target = r + \gamma Q'(s', \arg \max_{a'}(Q(s', a')))$   
        **end if**  
  
        Update the Q-network weights using gradient descent:  
         $L(\theta) = \text{MSE}(Q(s, a), target)$   
         $\theta \leftarrow \theta - \alpha \nabla_{\theta} L(\theta)$   
        Backpropagate the loss and update the Q-network weights  
        Update state  $s = s'$   
    **end for**  
    Reduce epsilon (exploration rate) over time (e.g.,  $\epsilon* = \epsilon\_decay$ )  
    **if** episode%target\_update\_frequency == 0 **then**  
        Update the target Q-network  $Q'$  with the current Q-network  $Q$  weights  
    **end if**  
    **if** episode%save\_model\_frequency == 0 **then**  
        Save the Q-network  $Q$  model  
    **end if**  
**end for**



# Improvements - Injecting experience

Before start: inject inside the agent's replay memory buffer some experiences derived from some episodes played by an human agent



# Results and metrics

- Results reflects more or less what we could expect:
  - DDQN variants achieved highest scores
  - Variants with injected experience seem to perform better



**Thank you for your attention!**